



**Московский государственный технический университет
им. Н.Э. Баумана
(МГТУ им. Н.Э. Баумана)
Радиотехнический факультет (РТ)**

**Отчёт по лабораторным работам №4-5
По дисциплине
«Технологии машинного обучения»**

Проверил:

Преподаватель кафедры ИУ-5

Гапанюк Ю.Е.

Подпись: _____

«__» _____ 2020 г.

Выполнил:

студент группы РТ5-61Б

Ануров Н.С.

Подпись: _____

«__» _____ 2020 г.

Москва, 2020

Цель лабораторной работы: изучение сложных способов подготовки выборки и подбора гиперпараметров на примере метода ближайших соседей.

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.neighbors import KNeighborsClassifier
from sklearn.model_selection import cross_val_score, train_test_split
from sklearn.preprocessing import LabelEncoder
pd.options.mode.chained_assignment = None
```

```
In [2]: df=pd.read_csv('weatherAUS.csv')
```

```
In [3]: df.head()
```

Out[3]:

	Date	Location	MinTemp	MaxTemp	Rainfall	Evaporation	Sunshine	WindGustDir	WindGustSpeed	WindDir9am	...	Humidity3pm	Pressure9am	Pressure
0	2008-12-01	Albury	13.4	22.9	0.6	NaN	NaN	W	44.0	W	...	22.0	1007.7	10
1	2008-12-02	Albury	7.4	25.1	0.0	NaN	NaN	WNW	44.0	NNW	...	25.0	1010.6	10
2	2008-12-03	Albury	12.9	25.7	0.0	NaN	NaN	WSW	46.0	W	...	30.0	1007.6	10
3	2008-12-04	Albury	9.2	28.0	0.0	NaN	NaN	NE	24.0	SE	...	16.0	1017.6	10
4	2008-12-05	Albury	17.5	32.3	1.0	NaN	NaN	W	41.0	ENE	...	33.0	1010.8	10

5 rows × 24 columns

```
In [4]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 142193 entries, 0 to 142192
Data columns (total 24 columns):
#   Column              Non-Null Count  Dtype  
---  -
0   Date                 142193 non-null object  
1   Location             142193 non-null object  
2   MinTemp              141556 non-null float64 
3   MaxTemp              141871 non-null float64 
4   Rainfall             140787 non-null float64 
5   Evaporation          81350 non-null float64 
6   Sunshine             74377 non-null float64 
7   WindGustDir          132863 non-null object  
8   WindGustSpeed        132923 non-null float64 
9   WindDir9am           132180 non-null object  
10  WindDir3pm           138415 non-null object  
11  WindSpeed9am         140845 non-null float64 
12  WindSpeed3pm         139563 non-null float64 
13  Humidity9am          140419 non-null float64 
14  Humidity3pm          138583 non-null float64 
15  Pressure9am          128179 non-null float64 
16  Pressure3pm          128212 non-null float64 
17  Cloud9am             88536 non-null float64 
18  Cloud3pm             85099 non-null float64 
19  Temp9am              141289 non-null float64 
20  Temp3pm              139467 non-null float64 
21  RainToday            140787 non-null object  
22  RISK_MM              142193 non-null float64 
23  RainTomorrow         142193 non-null object  
dtypes: float64(17), object(7)
memory usage: 26.0+ MB
```

```
In [5]: dt=df.drop(['Date','Location','WindGustDir','WindDir9am','WindDir3pm','RainToday'],axis=1)
dt.head()
```

Out[5]:

	MinTemp	MaxTemp	Rainfall	Evaporation	Sunshine	WindGustSpeed	WindSpeed9am	WindSpeed3pm	Humidity9am	Humidity3pm	Pressure9am	Pressure
0	13.4	22.9	0.6	NaN	NaN	44.0	20.0	24.0	71.0	22.0	1007.7	10
1	7.4	25.1	0.0	NaN	NaN	44.0	4.0	22.0	44.0	25.0	1010.6	10
2	12.9	25.7	0.0	NaN	NaN	46.0	19.0	26.0	38.0	30.0	1007.6	10
3	9.2	28.0	0.0	NaN	NaN	24.0	11.0	9.0	45.0	16.0	1017.6	10
4	17.5	32.3	1.0	NaN	NaN	41.0	7.0	20.0	82.0	33.0	1010.8	10

```
In [6]: dt['RainTomorrow'].value_counts()
```

```
Out[6]: No      110316
        Yes      31877
        Name: RainTomorrow, dtype: int64
```

```
In [7]: df_rain=pd.get_dummies(dt)
        df_rain.columns
```

```
Out[7]: Index(['MinTemp', 'MaxTemp', 'Rainfall', 'Evaporation', 'Sunshine',
              'WindGustSpeed', 'WindSpeed9am', 'WindSpeed3pm', 'Humidity9am',
              'Humidity3pm', 'Pressure9am', 'Pressure3pm', 'Cloud9am', 'Cloud3pm',
              'Temp9am', 'Temp3pm', 'RISK_MM', 'RainTomorrow_No', 'RainTomorrow_Yes'],
              dtype='object')
```

```
In [8]: df_rain.head()
```

```
Out[8]:
```

	MinTemp	MaxTemp	Rainfall	Evaporation	Sunshine	WindGustSpeed	WindSpeed9am	WindSpeed3pm	Humidity9am	Humidity3pm	Pressure9am	Pressure3pm
0	13.4	22.9	0.6	NaN	NaN	44.0	20.0	24.0	71.0	22.0	1007.7	1007.7
1	7.4	25.1	0.0	NaN	NaN	44.0	4.0	22.0	44.0	25.0	1010.6	1010.6
2	12.9	25.7	0.0	NaN	NaN	46.0	19.0	26.0	38.0	30.0	1007.6	1007.6
3	9.2	28.0	0.0	NaN	NaN	24.0	11.0	9.0	45.0	16.0	1017.6	1017.6
4	17.5	32.3	1.0	NaN	NaN	41.0	7.0	20.0	82.0	33.0	1010.8	1010.8

```
In [9]: df_rain['RainTomorrow_Yes'].value_counts()
```


```
Out[9]: 0      110316
        1       31877
        Name: RainTomorrow_Yes, dtype: int64
```

```
In [10]: data=df_rain.fillna(df_rain.mean())
         data
```

```
Out[10]:
```

	MinTemp	MaxTemp	Rainfall	Evaporation	Sunshine	WindGustSpeed	WindSpeed9am	WindSpeed3pm	Humidity9am	Humidity3pm	Pressure9am	Pressure3pm
0	13.4	22.9	0.6	5.469824	7.624853	44.0	20.0	24.0	71.0	22.0	1007.7	1007.7
1	7.4	25.1	0.0	5.469824	7.624853	44.0	4.0	22.0	44.0	25.0	1010.6	1010.6
2	12.9	25.7	0.0	5.469824	7.624853	46.0	19.0	26.0	38.0	30.0	1007.6	1007.6
3	9.2	28.0	0.0	5.469824	7.624853	24.0	11.0	9.0	45.0	16.0	1017.6	1017.6
4	17.5	32.3	1.0	5.469824	7.624853	41.0	7.0	20.0	82.0	33.0	1010.8	1010.8
...
142188	3.5	21.8	0.0	5.469824	7.624853	31.0	15.0	13.0	59.0	27.0	1024.7	1024.7
142189	2.8	23.4	0.0	5.469824	7.624853	31.0	13.0	11.0	51.0	24.0	1024.6	1024.6
142190	3.6	25.3	0.0	5.469824	7.624853	22.0	13.0	9.0	56.0	21.0	1023.5	1023.5
142191	5.4	26.9	0.0	5.469824	7.624853	37.0	9.0	9.0	53.0	24.0	1021.0	1021.0

```
In [12]: plt.scatter( x='RISK_MM', y='RainTomorrow_Yes', data=data);
```



```
In [13]: y=data['RainTomorrow_Yes'].values
X=data.drop(['RainTomorrow_Yes'],axis=1).values
X
```

```
Out[13]: array([[13.4, 22.9,  0.6, ..., 21.8,  0. ,  1. ],
 [ 7.4, 25.1,  0. , ..., 24.3,  0. ,  1. ],
 [12.9, 25.7,  0. , ..., 23.2,  0. ,  1. ],
 ...,
 [ 3.6, 25.3,  0. , ..., 24.5,  0. ,  1. ],
 [ 5.4, 26.9,  0. , ..., 26.1,  0. ,  1. ],
 [ 7.8, 27. ,  0. , ..., 26. ,  0. ,  1. ]])
```

```
In [14]: X_train,X_test,y_train,y_test=train_test_split(X,y,test_size=0.3,random_state=21)
np.unique(X_train)
```

```
Out[14]: array([ -8.5,  -8.2,  -8. , ..., 1040.4, 1040.6, 1040.9])
```

```
In [15]: np.unique(X_test)
```

```
Out[15]: array([ -7.5,  -7.2,  -7. , ..., 1040.4, 1040.5, 1041. ])
```

```
In [16]: knn=KNeighborsClassifier(n_neighbors=5)
```

```
In [17]: knn.fit(X_train,y_train)
```

```
Out[17]: KNeighborsClassifier(algorithm='auto', leaf_size=30, metric='minkowski',
                             metric_params=None, n_jobs=None, n_neighbors=5, p=2,
                             weights='uniform')
```

```
In [18]: knn.score(X_test,y_test)
```

```
Out[18]: 0.892892306249707
```

```

In [19]: cross_val_score(knn,X_train,y_train,cv=5)
Out[19]: array([0.89265083, 0.89194756, 0.89877932, 0.89164615, 0.89099312])

In [20]: np.mean(cross_val_score(knn,X_train,y_train,cv=5))
Out[20]: 0.8932033957904254

In [21]: from sklearn.model_selection import GridSearchCV

In [22]: knn_params = {'n_neighbors' : list(range(1,15))}

In [23]: knn_grid = GridSearchCV(knn,knn_params,cv=5)

In [24]: knn_grid.fit(X_train,y_train)
Out[24]: GridSearchCV(cv=5, error_score=nan,
                      estimator=KNeighborsClassifier(algorithm='auto', leaf_size=30,
                                                    metric='minkowski',
                                                    metric_params=None, n_jobs=None,
                                                    n_neighbors=5, p=2,
                                                    weights='uniform'),
                      iid='deprecated', n_jobs=None,
                      param_grid={'n_neighbors': [1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12,
                                                    13, 14]},
                      pre_dispatch='2*n_jobs', refit=True, return_train_score=False,
                      scoring=None, verbose=0)

In [25]: knn_grid.best_score_, knn_grid.best_params_
Out[25]: (0.8940473200381774, {'n_neighbors': 7})

```

Лабораторная работа №5

Цель лабораторной работы: изучение линейных моделей, SVM и деревьев решений.

```

In [26]: from sklearn.tree import DecisionTreeClassifier
         from sklearn.svm import SVC

In [27]: first_tree = DecisionTreeClassifier(random_state=17)

In [28]: cross_val_score(first_tree, X_train, y_train, cv=5)
Out[28]: array([1., 1., 1., 1., 1.])

In [29]: np.mean(cross_val_score(first_tree, X_train, y_train, cv=5))
Out[29]: 1.0

```

```
In [30]: tree_params = {'max_depth': np.arange(2, 11), 'max_features': [.5, .7, 1]}
```

```
In [31]: tree_grid = GridSearchCV(first_tree, tree_params, cv=5, n_jobs=-1)
```

```
In [32]: tree_grid.fit(X_train, y_train)
```

```
Out[32]: GridSearchCV(cv=5, error_score=nan,
                    estimator=DecisionTreeClassifier(ccp_alpha=0.0, class_weight=None,
                                                    criterion='gini', max_depth=None,
                                                    max_features=None,
                                                    max_leaf_nodes=None,
                                                    min_impurity_decrease=0.0,
                                                    min_impurity_split=None,
                                                    min_samples_leaf=1,
                                                    min_samples_split=2,
                                                    min_weight_fraction_leaf=0.0,
                                                    presort='deprecated',
                                                    random_state=17,
                                                    splitter='best'),
                    iid='deprecated', n_jobs=-1,
                    param_grid={'max_depth': array([ 2,  3,  4,  5,  6,  7,  8,  9, 10]),
                                'max_features': [0.5, 0.7, 1]},
                    pre_dispatch='2*n_jobs', refit=True, return_train_score=False,
                    scoring=None, verbose=0)
```

```
In [33]: tree_grid.best_score_, tree_grid.best_params_
```

```
Out[33]: (1.0, {'max_depth': 2, 'max_features': 0.5})
```

```
In [34]: from sklearn.metrics import accuracy_score
```

```
In [35]: tree_test_pred = tree_grid.predict(X_test)
```

```
In [36]: accuracy_score(y_test, tree_test_pred)
```

```
Out[36]: 1.0
```

```
In [37]: from sklearn.tree import export_graphviz
```

```
In [38]: second_tree = DecisionTreeClassifier(max_depth=5).fit(X_train, y_train)
second_tree.score(X_test, y_test)
```

```
Out[38]: 1.0
```

```
In [39]: from sklearn.svm import SVC, NuSVC, LinearSVC, OneClassSVM, SVR, NuSVR, LinearSVR
```

```
In [40]: svc = SVC(gamma='auto')
svc.fit(X_train, y_train)
```

```
Out[40]: SVC(C=1.0, break_ties=False, cache_size=200, class_weight=None, coef0=0.0,
            decision_function_shape='ovr', degree=3, gamma='auto', kernel='rbf',
            max_iter=-1, probability=False, random_state=None, shrinking=True,
            tol=0.001, verbose=False)
```

```
In [41]: svc.fit(X_train, y_train)
```

```
Out[41]: SVC(C=1.0, break_ties=False, cache_size=200, class_weight=None, coef0=0.0,
            decision_function_shape='ovr', degree=3, gamma='auto', kernel='rbf',
            max_iter=-1, probability=False, random_state=None, shrinking=True,
            tol=0.001, verbose=False)
```

```
In [42]: svc_pred = svc.predict(X_test)
```

```
In [43]: cross_val_score(svc, X_train, y_train, cv=5)
```

```
Out[43]: array([0.77540564, 0.77545587, 0.77545587, 0.77555634, 0.77545587])
```

```
In [48]: np.mean([0.77540564, 0.77545587, 0.77545587, 0.77555634, 0.77545587])
```

```
Out[48]: 0.775465918
```

```
In [49]: accuracy_score(y_test, svc_pred)
```

```
Out[49]: 0.7771109756669323
```

```
In [50]: from sklearn.linear_model import LinearRegression, LogisticRegressionCV
```

```
In [52]: y_new = data['RainTomorrow_Yes'].values
X_new = data.drop('RainTomorrow_Yes', axis=1)
```

```
In [53]: X_new_train, X_new_test, y_new_train, y_new_test = train_test_split(X_new, y_new, test_size=0.2, random_state=42)
```

```
In [54]: reg = LinearRegression()
```

```
In [55]: reg.fit(X_new_train, y_new_train)
```

```
Out[55]: LinearRegression(copy_X=True, fit_intercept=True, n_jobs=None, normalize=False)
```

```
In [56]: from sklearn.model_selection import StratifiedKFold
skf = StratifiedKFold(n_splits=5, shuffle=True, random_state=17)

c_values = np.logspace(-2, 3, 500)

logit_searcher = LogisticRegressionCV(Cs=c_values, cv=skf, verbose=1, n_jobs=-1)
logit_searcher.fit(X_new, y_new)

[Parallel(n_jobs=-1)]: Using backend LokyBackend with 8 concurrent workers.
[Parallel(n_jobs=-1)]: Done 2 out of 5 | elapsed: 9.6min remaining: 14.4min
[Parallel(n_jobs=-1)]: Done 5 out of 5 | elapsed: 10.1min finished
C:\Users\Nikita Anurov\anaconda3\lib\site-packages\sklearn\linear_model\_logistic.py:940: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:
https://scikit-learn.org/stable/modules/preprocessing.html
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear\_model.html#logistic-regression
extra_warning_msg=_LOGISTIC_SOLVER_CONVERGENCE_MSG)
```

```
Out[56]: LogisticRegressionCV(Cs=array([1.00000000e-02, 1.02334021e-02, 1.04722519e-02, 1.07166765e-02,
1.09668060e-02, 1.12227736e-02, 1.14847155e-02, 1.17527712e-02,
1.20270833e-02, 1.23077980e-02, 1.25950646e-02, 1.28890361e-02,
1.31898690e-02, 1.34977233e-02, 1.38127630e-02, 1.41351558e-02,
1.44650734e-02, 1.48026913e-02, 1.51481892e-02, 1.55017512e-02,
1.58635653e-02, 1.62...
8.50863158e+02, 8.70722485e+02, 8.91045332e+02, 9.11842520e+02,
9.33125118e+02, 9.54904456e+02, 9.77192128e+02, 1.00000000e+03]),
class_weight=None,
cv=StratifiedKFold(n_splits=5, random_state=17, shuffle=True),
dual=False, fit_intercept=True, intercept_scaling=1.0,
l1_ratios=None, max_iter=100, multi_class='auto',
n_jobs=-1, penalty='l2', random_state=None, refit=True,
scoring=None, solver='lbfgs', tol=0.0001, verbose=1)
```

```
In [57]: plt.plot(c_values, np.mean(logit_searcher.scores_[1], axis=0))
plt.xlabel('C')
plt.ylabel('Mean CV-accuracy');
```

